

Lecture 5: ML Frameworks

CS4787/5777 — Principles of Large-Scale Machine Learning Systems

Continuing from last time: Reverse-Mode AD

- Fix one output ℓ over \mathbb{R}
- Compute partial derivatives $\frac{\partial \ell}{\partial y}$ for each value y
- Need to do this by going *backward* through the computation

"Deep Learning" ML Frameworks

Classical Core Components:

- Numerical linear algebra library
- Hardware support (e.g. GPU)
- Backpropagation engine
- Library for expressing deep neural networks

All embedded in a high-level language

- Usually python.

Numerical Linear Algebra

You've already seen and used this sort of thing: NumPy.

- Arrays are objects "owned" by the library
- Any arithmetic operation on these objects goes through the library
 - The library calls an optimized function to compute the operation
 - This happens outside the python interpreter
 - Control is returned to python when the function finishes
 - By default you're only going to be running one such function at a time.

Numerical Linear Algebra: More Details

- Arrays are mutable
- Multiple references can exist!

Numerical Linear Algebra On-Device

- The simplest version of this is essentially a "copy" of NumPy for each sort of hardware we want to run on. This contains a copy of every function we want to support for each type of hardware.
 - e.g. one copy that runs on the CPU, one copy that runs on the GPU
- Arrays are located explicitly on one device
 - in PyTorch, you move them with `x.to("device_name")`
- When we try to call a function, the library checks where the inputs are located
 - if they're all on one device, it calls that device's version of the function
 - if they're not all on the same device, it raises an exception

Eager Execution vs Graph Execution

When we manifest a node of the compute graph, we can either:

- (eager) compute and manifest the value at that node immediately
- (graph) just manifest the node
 - need to call some function to compute the forward pass later

This was the classic distinction between TensorFlow and PyTorch

```
In [1]: import torch
```

```
In [3]: x = torch.ones((1,3,4))
```

```
In [6]: x = torch.ones(())
```

```
In [7]: x
```

```
Out[7]: tensor(1.)
```

```
In [46]: x = torch.ones(())
x.requires_grad = True
u = (x + 2)
y = u.square() # (x + 2)^2 --> 2 * (1 + 2) = 6
```

Advantages of eager mode (compute values & manifest graph at the same time):

- much better for value-based debugging!
- varying shapes
- less complicated
- condition on values

Advantages of lazy mode/graph mode (manifest graph first, then compute values):

- heavier static optimization
- better for shape-based debugging

- could use less memory
- graph overhead less/amortized

In [65]: `import time`

In [86]: `N**2 / (1024**3)`

Out[86]: 0.25

In [84]: `N = 1024 * 16`
`X = torch.randn(N,N)`
`Y = torch.randn(N,N)`

In [90]: `begin = time.time()`
`Z = X + Y`
`print(Z[0,0])`
`end = time.time()`
`print(f"elapsed: {(end - begin) * 1000} ms")`

tensor(-1.4525)
 elapsed: 60.443878173828125 ms

In [91]: `X_mps = X.to("mps")`
`Y_mps = Y.to("mps")`
`begin = time.time()`
`Z_mps = X_mps + Y_mps`
`print(Z_mps[0,0])`
`end = time.time()`
`print(f"elapsed: {(end - begin) * 1000} ms")`

tensor(-1.4525, device='mps:0')
 elapsed: 56.93316459655762 ms

In [83]: `554.5499324798584 / 118.25203895568848`

Out[83]: 4.689559160055244

In [89]: `4210.312843322754 / 985.3289127349854`

Out[89]: 4.273002434929221

In [98]: `U = torch.ones((1))`

In [99]: `U_mps = U.to("mps")`

In [100]: `U`

Out[100]: tensor([1.])

In [101]: `U_mps[0] = 3`

In [102]: `U`

Out[102]: tensor([1.])

```
In [103... U_mps
```

```
Out[103]: tensor([3.], device='mps:0')
```

```
In [104... U_mps.cpu()
```

```
Out[104]: tensor([3.])
```

```
In [105... U_mps.to("cpu")
```

```
Out[105]: tensor([3.])
```

```
In [106... U = torch.ones((5,5,5),device="mps")
```

```
In [111... (N ** 3)/(1024**3)
```

```
Out[111]: 4096.0
```

```
In [113... U = torch.ones((N,N,N),device="meta")
```

```
In [115... torch.ones((10000,20000),device="meta") @ torch.ones((20000,30000),device="meta")
```

```
Out[115]: tensor(..., device='meta', size=(10000, 30000))
```

```
In [116... torch.ones((10000,20000),device="meta") @ torch.ones((30000,30000),device="meta")
```

```
-----
RuntimeError                                Traceback (most recent call last)
Cell In[116], line 1
----> 1 torch.ones((10000,20000),device="meta") @ torch.ones((30000,30000),device="meta")

File /opt/anaconda3/lib/python3.9/site-packages/torch/_meta_registrations.py:448, in meta_mm(a, b)
    446 N, M1 = a.shape
    447 M2, P = b.shape
--> 448 check(M1 == M2, lambda: "a and b must have same reduction dim")
    449 return a.new_empty(N, P)

File /opt/anaconda3/lib/python3.9/site-packages/torch/_prims_common/__init__.py:1563, in check(b, s, exc_type)
    1556 """
    1557 Helper function for raising an error_type (default: RuntimeError) if a
boolean condition fails.
    1558 Error message is a callable producing a string (to avoid wasting time
    1559 string formatting in non-error case, and also to make it easier for to
rchedynamo
    1560 to trace.)
    1561 """
    1562 if not b:
-> 1563     raise exc_type(s())

RuntimeError: a and b must have same reduction dim
```

```
In [119... torch.ones((1000000,2000000),device="meta")
```

```
Out[119]: tensor(..., device='meta', size=(1000000, 2000000))
```

```
In [ ]: torch.ones((1000000,2000000)).to("meta")
```

```
In [1]: import torch
```

```
In [3]: X = torch.ones((1000,2000),device="meta")
```

```
In [4]: X.nonzero()
```

NotImplementedError

Traceback (most recent call last)

Cell In[4], line 1

----> 1 X.nonzero()

NotImplementedError: Could not run 'aten::nonzero' with arguments from the 'Meta' backend. This could be because the operator doesn't exist for this backend, or was omitted during the selective/custom build process (if using custom build). If you are a Facebook employee using PyTorch on mobile, please visit <https://fburl.com/ptmfixes> for possible resolutions. 'aten::nonzero' is only available for these backends: [CPU, MPS, BackendSelect, Python, FuncTorchDynamicLayerBackMode, Functionalize, Named, Conjugate, Negative, ZeroTensor, ADInplaceOrView, AutogradOther, AutogradCPU, AutogradCUDA, AutogradHIP, AutogradXLA, AutogradMPS, AutogradIPU, AutogradXPU, AutogradHPU, AutogradVE, AutogradLazy, AutogradMeta, AutogradMTIA, AutogradPrivateUse1, AutogradPrivateUse2, AutogradPrivateUse3, AutogradNestedTensor, Tracer, AutocastCPU, AutocastCUDA, FuncTorchBatched, FuncTorchVmapMode, Batched, VmapMode, FuncTorchGradWrapper, PythonTLSSnapshot, FuncTorchDynamicLayerFrontMode, PythonDispatcher].

CPU: registered at /Users/runner/work/pytorch/pytorch/pytorch/build/aten/src/ATen/RegisterCPU.cpp:31034 [kernel]

MPS: registered at /Users/runner/work/pytorch/pytorch/pytorch/build/aten/src/ATen/RegisterMPS.cpp:22748 [kernel]

BackendSelect: fallthrough registered at /Users/runner/work/pytorch/pytorch/pytorch/aten/src/ATen/core/BackendSelectFallbackKernel.cpp:3 [backend fallback]

Python: registered at /Users/runner/work/pytorch/pytorch/pytorch/aten/src/ATen/core/PythonFallbackKernel.cpp:144 [backend fallback]

FuncTorchDynamicLayerBackMode: registered at /Users/runner/work/pytorch/pytorch/pytorch/aten/src/ATen/funcTorch/DynamicLayer.cpp:491 [backend fallback]

Functionalize: registered at /Users/runner/work/pytorch/pytorch/pytorch/aten/src/ATen/FunctionalizeFallbackKernel.cpp:280 [backend fallback]

Named: registered at /Users/runner/work/pytorch/pytorch/pytorch/aten/src/ATen/core/NamedRegistrations.cpp:7 [backend fallback]

Conjugate: registered at /Users/runner/work/pytorch/pytorch/pytorch/aten/src/ATen/ConjugateFallback.cpp:17 [backend fallback]

Negative: registered at /Users/runner/work/pytorch/pytorch/pytorch/aten/src/ATen/native/NegateFallback.cpp:19 [backend fallback]

ZeroTensor: registered at /Users/runner/work/pytorch/pytorch/pytorch/aten/src/ATen/ZeroTensorFallback.cpp:86 [backend fallback]

ADInplaceOrView: fallthrough registered at /Users/runner/work/pytorch/pytorch/pytorch/aten/src/ATen/core/VariableFallbackKernel.cpp:63 [backend fallback]

AutogradOther: registered at /Users/runner/work/pytorch/pytorch/pytorch/torch/csrc/autograd/generated/VariableType_0.cpp:15256 [autograd kernel]

AutogradCPU: registered at /Users/runner/work/pytorch/pytorch/pytorch/torch/csrc/autograd/generated/VariableType_0.cpp:15256 [autograd kernel]

AutogradCUDA: registered at /Users/runner/work/pytorch/pytorch/pytorch/torch/csrc/autograd/generated/VariableType_0.cpp:15256 [autograd kernel]

AutogradHIP: registered at /Users/runner/work/pytorch/pytorch/pytorch/torch/csrc/autograd/generated/VariableType_0.cpp:15256 [autograd kernel]

AutogradXLA: registered at /Users/runner/work/pytorch/pytorch/pytorch/torch/csrc/autograd/generated/VariableType_0.cpp:15256 [autograd kernel]

AutogradMPS: registered at /Users/runner/work/pytorch/pytorch/pytorch/torch/csrc/autograd/generated/VariableType_0.cpp:15256 [autograd kernel]

AutogradIPU: registered at /Users/runner/work/pytorch/pytorch/pytorch/torch/csrc/autograd/generated/VariableType_0.cpp:15256 [autograd kernel]

AutogradXPU: registered at /Users/runner/work/pytorch/pytorch/pytorch/torch/csrc/autograd/generated/VariableType_0.cpp:15256 [autograd kernel]

AutogradHPU: registered at /Users/runner/work/pytorch/pytorch/pytorch/torch/csrc/autograd/generated/VariableType_0.cpp:15256 [autograd kernel]

AutogradVE: registered at /Users/runner/work/pytorch/pytorch/pytorch/torch/csrc

```
c/autograd/generated/VariableType_0.cpp:15256 [autograd kernel]
AutogradLazy: registered at /Users/runner/work/pytorch/pytorch/pytorch/torch/c
src/autograd/generated/VariableType_0.cpp:15256 [autograd kernel]
AutogradMeta: registered at /Users/runner/work/pytorch/pytorch/pytorch/torch/c
src/autograd/generated/VariableType_0.cpp:15256 [autograd kernel]
AutogradMTIA: registered at /Users/runner/work/pytorch/pytorch/pytorch/torch/c
src/autograd/generated/VariableType_0.cpp:15256 [autograd kernel]
AutogradPrivateUse1: registered at /Users/runner/work/pytorch/pytorch/pytorch/
torch/csrc/autograd/generated/VariableType_0.cpp:15256 [autograd kernel]
AutogradPrivateUse2: registered at /Users/runner/work/pytorch/pytorch/pytorch/
torch/csrc/autograd/generated/VariableType_0.cpp:15256 [autograd kernel]
AutogradPrivateUse3: registered at /Users/runner/work/pytorch/pytorch/pytorch/
torch/csrc/autograd/generated/VariableType_0.cpp:15256 [autograd kernel]
AutogradNestedTensor: registered at /Users/runner/work/pytorch/pytorch/pytorc
h/torch/csrc/autograd/generated/VariableType_0.cpp:15256 [autograd kernel]
Tracer: registered at /Users/runner/work/pytorch/pytorch/pytorch/torch/csrc/au
tograd/generated/TraceType_0.cpp:16728 [kernel]
AutocastCPU: fallthrough registered at /Users/runner/work/pytorch/pytorch/pyto
rch/aten/src/ATen/autocast_mode.cpp:487 [backend fallback]
AutocastCUDA: fallthrough registered at /Users/runner/work/pytorch/pytorch/pyt
orch/aten/src/ATen/autocast_mode.cpp:354 [backend fallback]
FuncTorchBatched: registered at /Users/runner/work/pytorch/pytorch/pytorch/ate
n/src/ATen/functorch/BatchRulesDynamic.cpp:64 [kernel]
FuncTorchVmapMode: fallthrough registered at /Users/runner/work/pytorch/pytorc
h/pytorch/aten/src/ATen/functorch/VmapModeRegistrations.cpp:28 [backend fallba
ck]
Batched: registered at /Users/runner/work/pytorch/pytorch/pytorch/aten/src/ATe
n/LegacyBatchingRegistrations.cpp:1073 [backend fallback]
VmapMode: fallthrough registered at /Users/runner/work/pytorch/pytorch/pytorc
h/aten/src/ATen/VmapModeRegistrations.cpp:33 [backend fallback]
FuncTorchGradWrapper: registered at /Users/runner/work/pytorch/pytorch/pytorc
h/aten/src/ATen/functorch/TensorWrapper.cpp:210 [backend fallback]
PythonTLSSnapshot: registered at /Users/runner/work/pytorch/pytorch/pytorch/at
en/src/ATen/core/PythonFallbackKernel.cpp:152 [backend fallback]
FuncTorchDynamicLayerFrontMode: registered at /Users/runner/work/pytorch/pytor
ch/pytorch/aten/src/ATen/functorch/DynamicLayer.cpp:487 [backend fallback]
PythonDispatcher: registered at /Users/runner/work/pytorch/pytorch/pytorch/ate
n/src/ATen/core/PythonFallbackKernel.cpp:148 [backend fallback]
```

In []: